

Import required libraries

```
In [1]: # List of all the libraries required in the code for easy reference
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

Load data

```
In [2]: # Load the required data
housing_df=pd.read_csv("BostonHousing.csv")
```

```
In [3]: # view the data
print(housing_df)
housing_df.shape
# the dataset has 506 cases (rows) and 14 attributes (columns)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	
	PTRATIO	LSTAT	MEDV	CAT.	MEDV						
0	15.3	4.98	24.0		0						
1	17.8	9.14	21.6		0						
2	17.8	4.03	34.7		1						
3	18.7	2.94	33.4		1						
4	18.7	5.33	36.2		1						
..						
501	21.0	9.67	22.4		0						
502	21.0	9.08	20.6		0						
503	21.0	5.64	23.9		0						
504	21.0	6.48	22.0		0						
505	21.0	7.88	11.9		0						

[506 rows x 14 columns]

Out[3]: (506, 14)

Rename Column CAT.MEDV

```
In [4]: # rename column CAT.MEDV to CAT_MEDV for better understanding
housing_df=housing_df.rename(columns={'CAT. MEDV':'CAT_MEDV'})
```

```
In [5]: # view the updated dataset
housing_df.head()
```

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	LSTAT	N
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	4.98	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	9.14	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	4.03	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	2.94	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	5.33	

Overall statistics for each column

```
In [6]: # Describe the overall dataset
housing_df.describe()
```

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000

Count of missing values

```
In [7]: # count of missing values in each column
missingvalues=housing_df.isnull().sum()
print(missingvalues)

# the sum of the missing values are zero indicating that the data has no mi
```

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
LSTAT     0
MEDV      0
CAT_MEDV  0
dtype: int64
```

Pie chart for AGE at different levels

```

In [8]: # Categorize 'AGE' into different levels
bins = [0, 25, 50, 75, 100]
labels = ['0-25%', '26-50%', '51-75%', '76-100%']

#copy the dataframe
housing_df_pie = housing_df.copy()

# create bins for the dataframe and create comparison dataframe
housing_df_pie['AGE_Category'] = pd.cut(housing_df_pie['AGE'], bins=bins, labels=labels)
comparison_df = housing_df_pie.groupby(['AGE_Category', 'CAT_MEDV']).size()


# plott the pie chart
colors = ['lightcoral', 'lightblue']
explode = (0.1, 0) # explode the 1st slice for emphasis
fig, axes = plt.subplots(nrows=2,ncols=2, figsize=(10,8))

for (row_index, row), ax in zip(comparison_df.iterrows(), axes.flatten()):
    ax.pie(row, labels=row.index, autopct='%1.1f%%', colors=colors, startangle=90)
    ax.set_title(f"Median Value above 30K for Age {row_index}")
    ax.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle

#plt.tight_layout()
fig.legend(row.index,loc="lower left", fontsize = 'large')
fig.text(0.5,0.01, "0 = Does not belong to the category (i.e below 30k)\n1 = Belongs to the category (i.e above 30k)")

```

Out[8]: Text(0.5, 0.01, '0 = Does not belong to the category (i.e below 30k)\n1 = Belongs to the category (i.e above 30k)')



Rows of outliers based on PTRATIO

```

In [9]: # find the outliers for PTRATIO column irrescpetive of the values in CAT_ME

# get the interquartile range
Q1 = housing_df['PTRATIO'].quantile(0.25)
Q3 = housing_df['PTRATIO'].quantile(0.75)
IQR = Q3 - Q1

# define bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
print("The lower bound is ", round(lower_bound,4), "\nThe upper bound is ",

# identifying the outliers
outliers = housing_df['PTRATIO'][(housing_df['PTRATIO'] < lower_bound) | (
if not outliers.empty:
    # all the values for the outlier value
    outlier_value = housing_df.loc[outliers.index]
    print("\nAll Outliers:")
    print(outlier_value)
    #print(outlier_value.count())

# the data records seem to be genuine outliers.

```

The lower bound is 13.2

The upper bound is 24.4

All Outliers:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
196	0.04011	80.0	1.52	0	0.404	7.287	34.1	7.3090	2	329	
197	0.04666	80.0	1.52	0	0.404	7.107	36.6	7.3090	2	329	
198	0.03768	80.0	1.52	0	0.404	7.274	38.3	7.3090	2	329	
257	0.61154	20.0	3.97	0	0.647	8.704	86.9	1.8010	5	264	
258	0.66351	20.0	3.97	0	0.647	7.333	100.0	1.8946	5	264	
259	0.65665	20.0	3.97	0	0.647	6.842	100.0	2.0107	5	264	
260	0.54011	20.0	3.97	0	0.647	7.203	81.8	2.1121	5	264	
261	0.53412	20.0	3.97	0	0.647	7.520	89.4	2.1398	5	264	
262	0.52014	20.0	3.97	0	0.647	8.398	91.5	2.2885	5	264	
263	0.82526	20.0	3.97	0	0.647	7.327	94.5	2.0788	5	264	
264	0.55007	20.0	3.97	0	0.647	7.206	91.6	1.9301	5	264	
265	0.76162	20.0	3.97	0	0.647	5.560	62.8	1.9865	5	264	
266	0.78570	20.0	3.97	0	0.647	7.014	84.6	2.1329	5	264	
267	0.57834	20.0	3.97	0	0.575	8.297	67.0	2.4216	5	264	
268	0.54050	20.0	3.97	0	0.575	7.470	52.6	2.8720	5	264	

	PTRATIO	LSTAT	MEDV	CAT_MEDV
196	12.6	4.08	33.3	1
197	12.6	8.61	30.3	1
198	12.6	6.62	34.6	1
257	13.0	5.12	50.0	1
258	13.0	7.79	36.0	1
259	13.0	6.90	30.1	1
260	13.0	9.59	33.8	1
261	13.0	7.26	43.1	1
262	13.0	5.91	48.8	1
263	13.0	11.25	31.0	1
264	13.0	8.10	36.5	1
265	13.0	10.45	22.8	0
266	13.0	14.79	30.7	1
267	13.0	7.44	50.0	1
268	13.0	3.16	43.5	1

Fix outliers

```
In [10]: # outliers account for 2.9% of the entire dataset
# we chose to omit these outliers.

housing_df_new = housing_df[~((housing_df['PTRATIO'] < lower_bound) | (housing_df['LSTAT'] > upper_bound))]
housing_df_new.shape
#housing_df_new.isna().sum()

# the new dataset has 491 cases (rows) and 14 attributes (columns)
```

Out[10]: (491, 14)

Statistics for each column

```
In [11]: # the mean, standard deviation, min, max, median, length, and missing values for each column

pd.DataFrame({'mean': housing_df_new.mean(),
              'median': housing_df_new.median(),
              'min': housing_df_new.min(),
              'max': housing_df_new.max(),
              'range': housing_df_new.max() - housing_df_new.min(),
              'Std. dev': housing_df_new.std(),
              'length': len(housing_df_new),
              'miss_val': housing_df_new.isnull().sum(),
              })
```

Out[11]:

	mean	median	min	max	range	Std. dev	length	miss
CRIM	3.708250	0.24522	0.00632	88.9762	88.96988	8.714712	491	
ZN	10.733198	0.00000	0.00000	100.0000	100.00000	23.011313	491	
INDUS	11.370692	9.90000	0.46000	27.7400	27.28000	6.828344	491	
CHAS	0.071283	0.00000	0.00000	1.0000	1.00000	0.257560	491	
NOX	0.553653	0.53200	0.38500	0.8710	0.48600	0.116288	491	
RM	6.251493	6.18500	3.56100	8.7800	5.21900	0.675457	491	
AGE	68.405703	77.00000	2.90000	100.0000	97.10000	28.277211	491	
DIS	3.814045	3.27210	1.12960	12.1265	10.99690	2.103519	491	
RAD	9.706721	5.00000	1.00000	24.0000	23.00000	8.789577	491	
TAX	412.246436	337.00000	187.00000	711.0000	524.00000	169.440997	491	
PTRATIO	18.624644	19.10000	13.60000	22.0000	8.40000	1.965453	491	
LSTAT	12.801181	11.65000	1.73000	37.9700	36.24000	7.181111	491	
MEDV	22.091853	20.90000	5.00000	50.0000	45.00000	8.868464	491	
CAT_MEDV	0.142566	0.00000	0.00000	1.0000	1.00000	0.349986	491	

Histogram for each quantitative variables


```
In [12]: # plot histogram for each variable in the new dataframe

# List of quantitative variable names
quantitative_vars = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'LSTAT', 'MEDV']

# set up subplots
fig, axes = plt.subplots(nrows=4, ncols=4, figsize=(15, 15))
fig.subplots_adjust(hspace=0.5)

# flatten the axes for easier iteration
axes = axes.flatten()

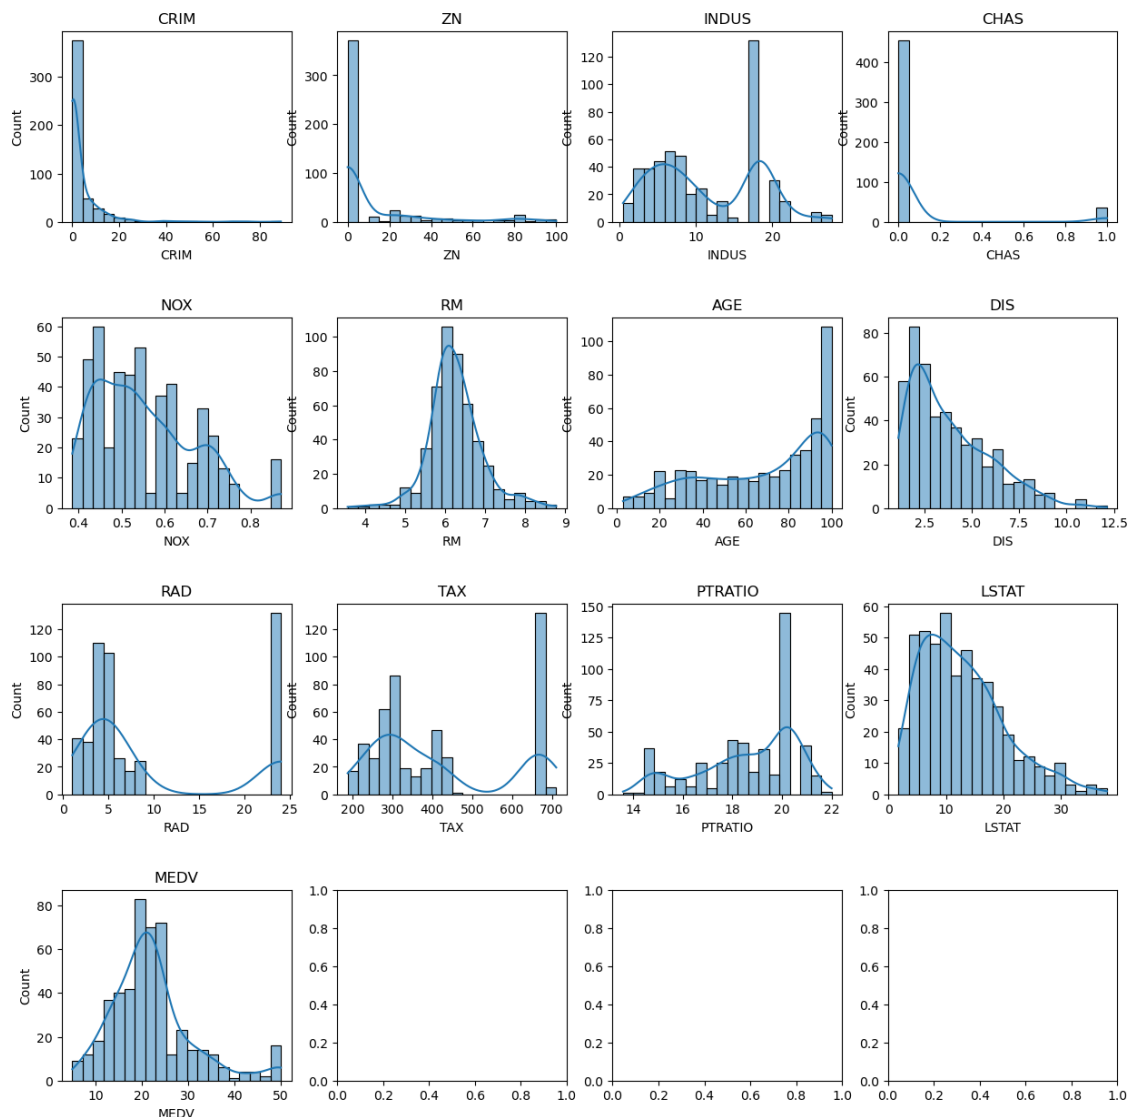
# plot histograms for each variable
for i, var in enumerate(quantitative_vars):
    sns.histplot(housing_df_new[var], ax=axes[i], bins=20, kde=True)
    axes[i].set_title(var)

plt.show()

# TAX variable has the highest level of variability.

# CRIM, ZN, CHAS, AGE, DIS, LSTAT are the variables which show skewness.

# the variables CRIM, AGE, DIS and LSTAT seem to have extreme values.
```



In [14]:

```

quantitative_vars = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS']

# Calculate the correlation matrix
correlation_matrix = housing_df_new[quantitative_vars].corr()
correlation_matrix_df = pd.DataFrame(housing_df_new[quantitative_vars].corr())

# Display the correlation table
print("Correlation Table:")
print(correlation_matrix_df)

```

Correlation Table:

	CRIM	ZN	INDUS	CHAS	NOX	RM	
AGE \							
CRIM	1.000000	-0.195751	0.402822	-0.059082	0.429326	-0.213881	0.358810
ZN	-0.195751	1.000000	-0.523917	-0.036383	-0.520639	0.297024	-0.576421
INDUS	0.402822	-0.523917	1.000000	0.054766	0.794809	-0.365529	0.668595
CHAS	-0.059082	-0.036383	0.054766	1.000000	0.094851	0.110079	0.089192
NOX	0.429326	-0.520639	0.794809	0.094851	1.000000	-0.338535	0.727605
RM	-0.213881	0.297024	-0.365529	0.110079	-0.338535	1.000000	-0.270840
AGE	0.358810	-0.576421	0.668595	0.089192	0.727605	-0.270840	1.000000
DIS	-0.389106	0.671706	-0.740027	-0.103407	-0.763542	0.237893	-0.745027
RAD	0.623634	-0.301859	0.589127	-0.012382	0.624151	-0.193630	0.464296
TAX	0.581076	-0.309810	0.715600	-0.042537	0.693280	-0.271667	0.524742
PTRATIO	0.293197	-0.368373	0.337121	-0.159898	0.237800	-0.277370	0.311267
LSTAT	0.453143	-0.406844	0.597411	-0.060225	0.606360	-0.614479	0.615702
MEDV	-0.391368	0.352029	-0.462905	0.198515	-0.473486	0.664089	-0.414797


	DIS	RAD	TAX	PTRATIO	LSTAT	MEDV
CRIM	-0.389106	0.623634	0.581076	0.293197	0.453143	-0.391368
ZN	0.671706	-0.301859	-0.309810	-0.368373	-0.406844	0.352029
INDUS	-0.740027	0.589127	0.715600	0.337121	0.597411	-0.462905
CHAS	-0.103407	-0.012382	-0.042537	-0.159898	-0.060225	0.198515
NOX	-0.763542	0.624151	0.693280	0.237800	0.606360	-0.473486
RM	0.237893	-0.193630	-0.271667	-0.277370	-0.614479	0.664089
AGE	-0.745027	0.464296	0.524742	0.311267	0.615702	-0.414797
DIS	1.000000	-0.506837	-0.560267	-0.288003	-0.512839	0.288708
RAD	-0.506837	1.000000	0.910676	0.470248	0.483260	-0.374980
TAX	-0.560267	0.910676	1.000000	0.451943	0.538496	-0.457212
PTRATIO	-0.288003	0.470248	0.451943	1.000000	0.362028	-0.453797
LSTAT	-0.512839	0.483260	0.538496	0.362028	1.000000	-0.743389
MEDV	0.288708	-0.374980	-0.457212	-0.453797	-0.743389	1.000000

Heatmap

```
In [15]: # Generate a heatmap (matrix plot) for the correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", lin
plt.title('Correlation Heatmap for Quantitative Variables')
plt.show()

# the variables TAX and RAD are strongly positively correlated: 0.910676.

# we can reduce the number of variables by removing the variables which are
```



Normalize new data

```
In [16]: # Normalize the data using Z-score normalization
normalized_data = StandardScaler().fit_transform(housing_df_new) #scales th

# Create a DataFrame with normalized data
normalized_df = pd.DataFrame(normalized_data, columns=housing_df_new.columns

# Compute the correlation matrix for normalized data
normalized_correlation_matrix = normalized_df[quantitative_vars].corr()
normalized_correlation_matrix
```


Out[16]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS
CRIM	1.000000	-0.195751	0.402822	-0.059082	0.429326	-0.213881	0.358810	-0.389106
ZN	-0.195751	1.000000	-0.523917	-0.036383	-0.520639	0.297024	-0.576421	0.671706
INDUS	0.402822	-0.523917	1.000000	0.054766	0.794809	-0.365529	0.668595	-0.740027
CHAS	-0.059082	-0.036383	0.054766	1.000000	0.094851	0.110079	0.089192	-0.103407
NOX	0.429326	-0.520639	0.794809	0.094851	1.000000	-0.338535	0.727605	-0.763542
RM	-0.213881	0.297024	-0.365529	0.110079	-0.338535	1.000000	-0.270840	0.237893
AGE	0.358810	-0.576421	0.668595	0.089192	0.727605	-0.270840	1.000000	-0.745027
DIS	-0.389106	0.671706	-0.740027	-0.103407	-0.763542	0.237893	-0.745027	1.000000
RAD	0.623634	-0.301859	0.589127	-0.012382	0.624151	-0.193630	0.464296	-0.506882
TAX	0.581076	-0.309810	0.715600	-0.042537	0.693280	-0.271667	0.524742	-0.560291
PTRATIO	0.293197	-0.368373	0.337121	-0.159898	0.237800	-0.277370	0.311267	-0.288032
LSTAT	0.453143	-0.406844	0.597411	-0.060225	0.606360	-0.614479	0.615702	-0.512819
MEDV	-0.391368	0.352029	-0.462905	0.198515	-0.473486	0.664089	-0.414797	0.288718

In [17]:

```
# Display the heatmap for normalized data
plt.figure(figsize=(10, 8))
sns.heatmap(normalized_correlation_matrix, annot=True, cmap='coolwarm', fmt='%.2f')
plt.title('Correlation Matrix Heatmap (Normalized Data)')
plt.show()

# the correlation values have not changed indicating that the variables in
```



Linear Regression

Split data in training and test dataset

```
In [18]: x = housing_df_new.drop('MEDV', axis =1)
y = housing_df_new['MEDV']

# split the dataset in training and test dataset, 80-20 split.
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, ra
```

Fit linear regression model

```
In [19]: model = LinearRegression()

# Fit the model on the training data
model.fit(x_train, y_train)
```

```
Out[19]: ▾ LinearRegression
LinearRegression()
```

Predict and Evaluate the model

```
In [20]: y_pred = model.predict(x_test)

# to understand how well the model is doing, we compare the MSE of the model
# calculate the mean value of the training dataset
mean_baseline = y_train.mean()

# create predictions based on the mean for the test set
y_pred_baseline = [mean_baseline] * len(y_test)

# evaluate the baseline model using MSE
mse_baseline = mean_squared_error([mean_baseline] * len(y_test), y_test)
print(f"Mean Squared Error (Baseline model): {mse_baseline}")

# evaluate the model (for example, using Mean Squared Error)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error:\t\t\t\t\t {mse}")

# this indicates the MSE value calculated is relatively low.
```

```
Mean Squared Error (Baseline model): 55.92864108270859
Mean Squared Error: 12.104631354705866
```

Extract the important features

```
In [21]: # get the coefficients from the model
coefficients = model.coef_

# create a DataFrame to display feature names and their corresponding coefficients
feature_importance = pd.DataFrame({
    'Feature': x.columns,
    'Coefficient': coefficients
})

# display the DataFrame
print(feature_importance)

# we ignore the coefficient for CAT_MEDV as the CAT_MEDV variable is derived
# The NOX has the highest negative coefficient, this states that an increase in NOX leads to a decrease in MEDV
# The second variable CHAS has a positive coefficient of 2.79 indicating that houses with a view of the Charles River have a higher median value
```

	Feature	Coefficient
0	CRIM	-0.122795
1	ZN	-0.005215
2	INDUS	0.137449
3	CHAS	2.790639
4	NOX	-15.558904
5	RM	0.542332
6	AGE	0.000753
7	DIS	-0.607606
8	RAD	0.152872
9	TAX	-0.006184
10	PTRATIO	-0.551640
11	LSTAT	-0.533303
12	CAT_MEDV	12.696295

In []: